

SYSTEM PROGRAMMING

WEEK 3: FILE I/O

Seongjin Lee

Updated: 2016-09-21
02-fileio

insight@hanyang.ac.kr

<http://esos.hanyang.ac.kr>

Esos Lab. Hanyang University



Table of contents

1. File I/Os

1.1 Standard I/O

2. File Sharing



FILE I/OS

File Descriptors

A *file descriptor* (or *file handle*) is a small, non-negative integer which identifies a file to the kernel.



File Descriptors

A *file descriptor* (or *file handle*) is a small, non-negative integer which identifies a file to the kernel.


Traditionally, `stdin`, `stdout` and `stderr` are 0, 1 and 2 respectively.



File Descriptors

A *file descriptor* (or *file handle*) is a small, non-negative integer which identifies a file to the kernel.


Traditionally, `stdin`, `stdout` and `stderr` are 0, 1 and 2 respectively.



File Descriptors

A *file descriptor* (or *file handle*) is a small, non-negative integer which identifies a file to the kernel.

Traditionally, `stdin`, `stdout` and `stderr` are 0, 1 and 2 respectively.



Relying on “magic numbers” is BAD. Use `STDIN_FILENO`, `STDOUT_FILENO` and `STDERR_FILENO` defined in `<unistd.h>` or `stdin`, `stdout`, and `stderr` defined in `<stdio.h>`.



How many files can you open? (./codes/fdcount.c)

```
1  long open_max(void)
2  {
3      if (openmax == 0) {
4          errno = 0;
5          /* first time through */
6          if ((openmax = sysconf(_SC_OPEN_MAX)) < 0) {
7              if (errno == 0)
8                  openmax = OPEN_MAX_GUESS; /* it is indeterminate */
9              else
10                 fprintf(stderr, "sysconf error for _SC_OPEN_MAX");
11             }
12     }
13     return(openmax);
14 }
15
16 int main(){
17     printf("The number of file descriptors a process can open: %d\n", (int)RLIMIT_
18         NOFILE);
19     printf("The number of file descriptors an user can open: %ld\n", openmax );
20     return 0;
21 }
```

How to compile

```
$ cd codes; cc -Wall -g -o fdcount fdcount.c
```



FILE I/OS : STANDARD I/O

Basic File I/Os

There are five fundamental UNIX file I/O related functions:

- `open(2)`
- `close(2)`
- `lseek(2)`
- `read(2)`
- `write(2)`



open(2)

```
#include <fcntl.h>
int open(const char *path, int oflag, ... /* mode_t mode */ );
int openat(int fd, const char *path, int oflag, ... /* mode_t mode */ );
```

Both return: file descriptor if OK, 1 on error

The *path* parameter is the name of the file to open or create.

Options are specified by the *oflag*



Options are

oflag must be one (and only one) of:

- O_RDONLY
– Open for reading only
- O_WRONLY
– Open for writing only
- O_RDWR
– Open for reading and writing

and may be OR'd with any of these:

- O_APPEND – Append to end of file for each write
- O_CREAT – Create the file if it doesn't exist. Requires *mode* argument
- O_EXCL – Generate error if O_CREAT and file already exists. (atomic)
- O_TRUNC – If file exists and successfully open in O_WRONLY or O_RDWR, make length = 0
- O_NOCTTY – If pathname refers to a terminal device, do not allocate the device as a controlling terminal
- O_NONBLOCK – If pathname refers to a FIFO, block special, or char special, set nonblocking mode (open and I/O)
- O_SYNC – Each write waits for physical I/O to complete



openat(2)

openat(2) function is equivalent to the open() function except in the case where the path specifies a relative path in an atomic fashion

- O_EXEC – Open for execute only
- O_SEARCH – Open for search only (applies to directories)
- O_DIRECTORY – If path resolves to a non-directory file, fail and set errno to ENOTDIR.
- O_DSYNC – Wait for physical I/O for data, except file attributes
- O_RSYNC – Block read operations on any pending writes.
- O_PATH – Obtain a file descriptor purely for fd-level operations. (Linux >2.6.36 only)



FILE SHARING

File Sharing

Atomicity of the file fundamental file I/O functions

File sharing

manipulation of file descriptors

