

# SYSTEM PROGRAMMING

## WEEK 7: SIGNALS

---

Seongjin Lee

Updated: 2016-10-19

07\_signal

[insight@hanyang.ac.kr](mailto:insight@hanyang.ac.kr)

<http://esos.hanyang.ac.kr>

Esos Lab. Hanyang University



# Table of contents

1. Signal Concepts
2. The issues
3. Use cases of signal
4. Last Words



This chapter covers following items

- The concept
- Use cases of signal
- The problems of earlier implementations
- The correct ways



# SIGNAL CONCEPTS



# Signal Concepts

Every signal has a name that begins with SIG and it is assigned with a positive number defined in `<signal.h>`

- SIGABRT: generated when a process calls `abort` function
- SIGALRM: generated when a timer set by `alarm` function goes off
- Different versions of UNIX have different number of signals



# Signal Concepts

## Signal generating conditions

1. terminal generated signals (DELETE or ^C on many systems) causes the interrupt signal (SIGINT)



# Signal Concepts

## Signal generating conditions

1. terminal generated signals (DELETE or ^C on many systems) causes the interrupt signal (SIGINT)
2. Hardware exceptions generate signals
  - invalid memory reference (SIGSEGV)
  - I/O completed (SIGIO)
  - user disconnected from the system (SIGHUP)
  - detected by HW and the kernel is notified
  - kernel generates the appropriate signal for the process



# Signal Concepts

## Signal generating conditions

1. terminal generated signals (DELETE or ^C on many systems) causes the interrupt signal (SIGINT)
2. Hardware exceptions generate signals
  - invalid memory reference (SIGSEGV)
  - I/O completed (SIGIO)
  - user disconnected from the system (SIGHUP)
  - detected by HW and the kernel is notified
  - kernel generates the appropriate signal for the process
3. kill(2) function allows a process to send any signal to another process or group (have to be owner, or the superuser)





# Signal Concepts

## Signal generating conditions

1. terminal generated signals (DELETE or ^C on many systems) causes the interrupt signal (SIGINT)
2. Hardware exceptions generate signals
  - invalid memory reference (SIGSEGV)
  - I/O completed (SIGIO)
  - user disconnected from the system (SIGHUP)
  - detected by HW and the kernel is notified
  - kernel generates the appropriate signal for the process
3. kill(2) function allows a process to send any signal to another process or group (have to be owner, or the superuser)
4. kill(1) command sends signal to other process



# Signal Concepts

## Signal generating conditions

1. terminal generated signals (DELETE or ^C on many systems) causes the interrupt signal (SIGINT)
2. Hardware exceptions generate signals
  - invalid memory reference (SIGSEGV)
  - I/O completed (SIGIO)
  - user disconnected from the system (SIGHUP)
  - detected by HW and the kernel is notified
  - kernel generates the appropriate signal for the process
3. kill(2) function allows a process to send any signal to another process or group (have to be owner, or the superuser)
4. kill(1) command sends signal to other process
5. Software conditions can generate signals



# Signal Concepts cont'd

Signals are asynchronous events that occurs randomly

The process has to tell the kernel “if and when this signal occurs, do the following”



# Signal Concepts cont'd

Signals are asynchronous events that occurs randomly

The process has to tell the kernel “if and when this signal occurs, do the following”

We can tell the kernel to do one of three things

1. **Ignore the signal** following two can never be ignored: SIGKILL and SIGSTOP



# Signal Concepts cont'd

Signals are asynchronous events that occurs randomly

The process has to tell the kernel “if and when this signal occurs, do the following”

We can tell the kernel to do one of three things

1. **Ignore the signal** following two can never be ignored: SIGKILL and SIGSTOP
2. **Catch the signal** We tell the kernel to call a customized function whenever the signal occurs
  - if SIGCHLD signal is caught, it means child has terminated
  - signal catching function calls waitpid to fetch the child's process ID and termination status



# Signal Concepts cont'd

Signals are asynchronous events that occurs randomly

The process has to tell the kernel “if and when this signal occurs, do the following”

We can tell the kernel to do one of three things

1. **Ignore the signal** following two can never be ignored: SIGKILL and SIGSTOP
2. **Catch the signal** We tell the kernel to call a customized function whenever the signal occurs
  - if SIGCHLD signal is caught, it means child has terminated
  - signal catching function calls waitpid to fetch the child's process ID and termination status
3. **Use default action** every signal has a default action
  - the default action for most signals is to terminate the process>



# Signal Concepts cont'd

| Name      | Description                       | ISO C | SUS | FreeBSD<br>8.0 | Linux<br>3.2.0 | Mac OS X<br>10.6.8 | Solaris<br>10 | Default action   |
|-----------|-----------------------------------|-------|-----|----------------|----------------|--------------------|---------------|------------------|
| SIGABRT   | abnormal termination (abort)      | •     | •   | •              | •              | •                  | •             | terminate+core   |
| SIGALRM   | timer expired (alarm)             |       | •   | •              | •              | •                  | •             | terminate        |
| SIGBUS    | hardware fault                    |       | •   | •              | •              | •                  | •             | terminate+core   |
| SIGCANCEL | threads library internal use      |       |     |                |                |                    | •             | ignore           |
| SIGCHLD   | change in status of child         |       | •   | •              | •              | •                  | •             | ignore           |
| SIGCONT   | continue stopped process          |       | •   | •              | •              | •                  | •             | continue/ignore  |
| SIGEMT    | hardware fault                    |       |     | •              | •              | •                  | •             | terminate+core   |
| SIGFPE    | arithmetic exception              | •     | •   | •              | •              | •                  | •             | terminate+core   |
| SIGFREEZE | checkpoint freeze                 |       |     |                |                |                    | •             | ignore           |
| SIGHUP    | hangup                            |       | •   | •              | •              | •                  | •             | terminate        |
| SIGILL    | illegal instruction               | •     | •   | •              | •              | •                  | •             | terminate+core   |
| SIGINFO   | status request from keyboard      |       |     | •              | •              | •                  |               | ignore           |
| SIGINT    | terminal interrupt character      | •     | •   | •              | •              | •                  | •             | terminate        |
| SIGIO     | asynchronous I/O                  |       |     | •              | •              | •                  | •             | terminate/ignore |
| SIGIOT    | hardware fault                    |       |     | •              | •              | •                  | •             | terminate+core   |
| SIGJVM1   | Java virtual machine internal use |       |     |                |                |                    | •             | ignore           |
| SIGJVM2   | Java virtual machine internal use |       |     |                |                |                    | •             | ignore           |
| SIGKILL   | termination                       |       | •   | •              | •              | •                  | •             | terminate        |
| SIGLOST   | resource lost                     |       |     |                |                |                    | •             | terminate        |
| SIGLWP    | threads library internal use      |       |     | •              |                |                    | •             | terminate/ignore |
| SIGPIPE   | write to pipe with no readers     |       | •   | •              | •              | •                  | •             | terminate        |
| SIGPOLL   | pollable event (poll)             |       |     |                | •              |                    | •             | terminate        |
| SIGPROF   | profiling time alarm (setitimer)  |       |     | •              | •              | •                  | •             | terminate        |

Figure: UNIX System signals



# signal Function

```
#include <signal.h>
void (*signal(int signo, void (*func)(int)))(int);
// Returns: previous disposition of signal (see following) if OK, SIG_ERR on
            error
```

- signo is name of the signal from the Table 10.1
- func is
  - SIG\_IGN : to ignore the signal
  - SIG\_DFL : to use the default value
  - address of a function to be called when the signal occurs—they are called *signal handler* or *signal-catching function*





# Signal Exmple Code: codes/usr\_sig.c |

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <signal.h> // for signalling
6  #include <errno.h>
7
8  static void sig_usr(int); /* one handler for both signals */
9
10 int
11 main(void) {
12     if (signal(SIGUSR1, sig_usr) == SIG_ERR) {
13         fprintf(stderr, "Can't catch SIGUSR1: %s", strerror(errno));
14         exit(1);
15     }
16     if (signal(SIGUSR2, sig_usr) == SIG_ERR) {
17         fprintf(stderr, "Can't catch SIGUSR2: %s", strerror(errno));
18         exit(1);
19     }
20
21     if (signal(SIGHUP, sig_usr) == SIG_ERR) {
22         fprintf(stderr, "Can't catch SIGHUP: %s", strerror(errno));
```



# Signal Exmple Code: codes/usr\_sig.c II

```
23     exit(1);
24 }
25
26 for ( ; ; )
27     pause();
28 }
29
30 static void
31 sig_usr(int signo) { /* argument is signal number */
32     if (signo == SIGUSR1)
33         printf("received SIGUSR1\n");
34     else if (signo == SIGUSR2)
35         printf("received SIGUSR2\n");
36     else if (signo == SIGHUP)
37         printf("received SIGHUP\n");
38     else {
39         fprintf(stderr, "received signal: %d\n", signo);
40         exit(1);
41     }
42     return;
43 }
```



# Signal Exmaples

invoke the program in the background and use `kill(1)` to send signal

```
James@maker:codes$ ./usr\_sig &  
[2] 4987  
James@maker:codes$ kill -USR1 4987  
received SIGUSR1  
James@maker:codes$ kill -USR2 4987  
received SIGUSR2  
James@maker:codes$ kill -HUP 4987  
received SIGHUP  
James@maker:codes$ kill -INT 4987  
[2]+ Interrupt: 2 ./usr\_sig  
James@maker:codes$
```




# signal Function

## At process creation

- When a process calls `fork`, the child inherits the parents signal disposition
- Child starts off with copy of the parent's memory image
- the address of a signal-catching function has meaning in the child



## THE ISSUES



# Unreliable Signals

In earlier versions of UNIX system, signals were unreliable



# Unreliable Signals

In earlier versions of UNIX system, signals were unreliable

1. signals get lost: signal occurs and the process never know about it



# Unreliable Signals

In earlier versions of UNIX system, signals were unreliable

1. signals get lost: signal occurs and the process never know about it
2. window of time—after the signal has occurred, but before the call to signal in the signal handler—when the interrupt signal could occur another time. The second signal would cause the default action to occur (terminate the process)





# Unreliable Signals

In earlier versions of UNIX system, signals were unreliable

1. signals get lost: signal occurs and the process never know about it
2. window of time—after the signal has occurred, but before the call to signal in the signal handler—when the interrupt signal could occur another time. The second signal would cause the default action to occur (terminate the process)
3. little control over a signal: unable to turn a signal off when it didn't want the signal to occur. All it can do is to catch or ignore the signal.



# Interrupted System Calls cont'd

The slow system calls are those that can block forever

- reads (for pipes, terminal devices, and network devices) that can block the caller
- writes that can block the caller forever if the data can't be accepted immediately
- open on a certain file types (terminal device) that block the caller until some condition occurs
- the pause and wait function
- certain ioctl operations
- some of interprocess communication function



# Interrupted System Calls cont'd

The problem with interrupted system calls is that error returns must be explicit

```
again:
    if ((n = read(fd, buf, BUFSIZE)) < 0 ) {
        if (errno == EINTR)
            goto again; /* just an interrupted system call */
        /* handle other errors */
    }
```

The solution of 4.2BSD was to introduce the automatic restarting of `ioctl`, `read`, `readv`, `write`, `writen`, `wait`, and `waitpid`



## Reentrant Functions

Reentrant functions are guaranteed to be safe to call from within a signal handler. They are also called *async-signal safe*. As a general rule, when calling the reentrant functions from a signal handler, we should save and restore `errno`.



## Reentrant Functions

Reentrant functions are guaranteed to be safe to call from within a signal handler. They are also called *async-signal safe*. As a general rule, when calling the reentrant functions from a signal handler, we should save and restore `errno`.

### Scenario

1. while process is running, it catches a signal. process is temporarily interrupted by the signal handler



## Reentrant Functions

Reentrant functions are guaranteed to be safe to call from within a signal handler. They are also called *async-signal safe*. As a general rule, when calling the reentrant functions from a signal handler, we should save and restore `errno`.

### Scenario

1. while process is running, it catches a signal. process is temporarily interrupted by the signal handler
2. instructions in signal handler executes, then returns



## Reentrant Functions

Reentrant functions are guaranteed to be safe to call from within a signal handler. They are also called *async-signal safe*. As a general rule, when calling the reentrant functions from a signal handler, we should save and restore `errno`.

### Scenario

1. while process is running, it catches a signal. process is temporarily interrupted by the signal handler
2. instructions in signal handler executes, then returns
3. upon return of signal handler, the process continues to run



## Reentrant Functions

Reentrant functions are guaranteed to be safe to call from within a signal handler. They are also called *async-signal safe*. As a general rule, when calling the reentrant functions from a signal handler, we should save and restore `errno`.

### Scenario

1. while process is running, it catches a signal. process is temporarily interrupted by the signal handler
2. instructions in signal handler executes, then returns
3. upon return of signal handler, the process continues to run

The problem is that signal handler can't tell whether a process was in the middle of execution. *The result becomes unpredictable*





# Reentrant Functions cont'd

|               |             |                   |             |                  |
|---------------|-------------|-------------------|-------------|------------------|
| abort         | faccessat   | linkat            | select      | socketpair       |
| accept        | fchmod      | listen            | sem_post    | stat             |
| access        | fchmodat    | lseek             | send        | symlink          |
| aio_error     | fchown      | lstat             | sendmsg     | symlinkat        |
| aio_return    | fchownat    | mkdir             | sendto      | tcdrain          |
| aio_suspend   | fcntl       | mkdirat           | setgid      | tcflow           |
| alarm         | fdatasync   | mkfifo            | setpgid     | tcflush          |
| bind          | fexecve     | mkfifoat          | setsid      | tcgetattr        |
| cfgetispeed   | fork        | mknod             | setsockopt  | tcgetpgrp        |
| cfgetospeed   | fstat       | mknodat           | setuid      | tcsendbreak      |
| cfsetispeed   | fstatat     | open              | shutdown    | tcsetattr        |
| cfsetospeed   | fsync       | openat            | sigaction   | tcsetpgrp        |
| chdir         | ftruncate   | pause             | sigaddset   | time             |
| chmod         | futimens    | pipe              | sigdelset   | timer_getoverrun |
| chown         | getegid     | poll              | sigemptyset | timer_gettime    |
| clock_gettime | geteuid     | posix_trace_event | sigfillset  | timer_settime    |
| close         | getgid      | pselect           | sigismember | times            |
| connect       | getgroups   | raise             | signal      | umask            |
| creat         | getpeername | read              | sigpause    | uname            |
| dup           | getpgrp     | readlink          | sigpending  | unlink           |
| dup2          | getpid      | readlinkat        | sigprocmask | unlinkat         |
| execl         | getppid     | recv              | sigqueue    | utime            |
| execle        | getsockname | recvfrom          | sigset      | utimensat        |
| execv         | getsockopt  | recvmsg           | sigsuspend  | utimes           |
| execve        | getuid      | rename            | sleep       | wait             |
| _Exit         | kill        | renameat          | socketmark  | waitpid          |
| _exit         | link        | rmdir             | socket      | write            |

Figure: Reentrant functions that may be called from a signal handler



# Reliable-Signal Terminology and Semantics

- **a signal is *generated* (or sent) for a process** when the event that causes the signal occurs
  - The event can be hardware exception, software condition, a terminal-generated signal, or a call to the `kill` function
  - when the signal is generated, the kernel usually sets a flag of some form in the process table



# Reliable-Signal Terminology and Semantics

- **a signal is *generated* (or sent) for a process** when the event that causes the signal occurs
  - The event can be hardware exception, software condition, a terminal-generated signal, or a call to the `kill` function
  - when the signal is generated, the kernel usually sets a flag of some form in the process table
- **a signal is *delivered* to a process** when the action for a signal is taken
  - during the time between the generation of a signal and its delivery, the signal is said to be *pending*



# Reliable-Signal Terminology and Semantics

- **a signal is *generated* (or sent) for a process** when the event that causes the signal occurs
  - The event can be hardware exception, software condition, a terminal-generated signal, or a call to the `kill` function
  - when the signal is generated, the kernel usually sets a flag of some form in the process table
- **a signal is *delivered* to a process** when the action for a signal is taken
  - during the time between the generation of a signal and its delivery, the signal is said to be *pending*
- **a process has the option of *blocking* the delivery of a signal**



# Reliable-Signal Terminology and Semantics

- **a signal is *generated* (or sent) for a process** when the event that causes the signal occurs
  - The event can be hardware exception, software condition, a terminal-generated signal, or a call to the `kill` function
  - when the signal is generated, the kernel usually sets a flag of some form in the process table
- **a signal is *delivered* to a process** when the action for a signal is taken
  - during the time between the generation of a signal and its delivery, the signal is said to be *pending*
- **a process has the option of *blocking* the delivery of a signal**
- signals are *queued* when system delivers signals



# Reliable-Signal Terminology and Semantics

- **a signal is *generated* (or sent) for a process** when the event that causes the signal occurs
  - The event can be hardware exception, software condition, a terminal-generated signal, or a call to the `kill` function
  - when the signal is generated, the kernel usually sets a flag of some form in the process table
- **a signal is *delivered* to a process** when the action for a signal is taken
  - during the time between the generation of a signal and its delivery, the signal is said to be *pending*
- **a process has the option of *blocking* the delivery of a signal**
- signals are *queued* when system delivers signals
- POSIX.1 does not specify the order if more than one signal is ready to be delivered to a process



# Reliable-Signal Terminology and Semantics

- **a signal is *generated* (or sent) for a process** when the event that causes the signal occurs
  - The event can be hardware exception, software condition, a terminal-generated signal, or a call to the `kill` function
  - when the signal is generated, the kernel usually sets a flag of some form in the process table
- **a signal is *delivered* to a process** when the action for a signal is taken
  - during the time between the generation of a signal and its delivery, the signal is said to be *pending*
- **a process has the option of *blocking* the delivery of a signal**
- signals are *queued* when system delivers signals
- POSIX.1 does not specify the order if more than one signal is ready to be delivered to a process
- **each process has a *signal mask*** that defines the set of signals currently blocked for delivery to that process



# Reliable-Signal Terminology and Semantics

- a **signal is *generated* (or sent) for a process** when the event that causes the signal occurs
  - The event can be hardware exception, software condition, a terminal-generated signal, or a call to the `kill` function
  - when the signal is generated, the kernel usually sets a flag of some form in the process table
- a **signal is *delivered* to a process** when the action for a signal is taken
  - during the time between the generation of a signal and its delivery, the signal is said to be *pending*
- a **process has the option of *blocking* the delivery of a signal**
- signals are *queued* when system delivers signals
- POSIX.1 does not specify the order if more than one signal is ready to be delivered to a process
- **each process has a *signal mask*** that defines the set of signals currently blocked for delivery to that process
- POSIX.1 defines a data type `sigset_t` that holds a *signal set*





## USE CASES OF SIGNAL



# kill and raise Functions

kill function sends a signal to a process or a group of processes

raise function allows a process to send a signal to itself

```
#include <signal.h>
int kill(pid_t pid, int signo);
int raise(int signo);
// Both return: 0 if OK, 1 on error
```

```
raise(signo); == kill(getpid(), signo);
```

- $pid > 0$  sends signal to pid
- $pid == 0$  sends signal to all processes in process group ID of the sender
- $pid < 0$  The signal is sent to all process in process group ID of  $|pid|$
- $pid == -1$  send signal to all processes on the system for which the sender has permission to send the signal



# alarm Function

alarm function allows us to set a time that will expire at a specified time in the future. When the timer expires, the SIGALRM signal is generated

```
#include <unistd.h>
unsigned int alarm(unsigned int seconds);
// Returns: 0 or number of seconds until previously set alarm
```

- only one of alarm clocks per process
- if previously registered alarm clock has not expired, the number of seconds left for that alarm clock is returned as the value of this function
- if previous alarm is not expired and *seconds* value is 0, the previous alarm is canceled
- most processes using alarm, catches this signal



# pause Function

pause function suspends the calling process until a signal is caught

```
#include <unistd.h>
int pause(void);
// Returns: 1 with errno set to EINTR
```



# alarm and pause Functions Example

codes/sleep-pause.c

```
1  #include <signal.h>
2  #include <unistd.h>
3
4  static void
5  sig_alrm(int signo)
6  {
7      /* nothing to do, just return to wake up the pause */
8  }
9
10 unsigned int
11 sleep1(unsigned int seconds)
12 {
13     if (signal(SIGALRM, sig_alrm) == SIG_ERR)
14         return(seconds);
15     alarm(seconds); /* start the timer */
16     pause(); /* next caught signal wakes us up */
17     return(alarm(0)); /* turn off timer, return unslept time */
18 }
```



# alarm and pause Functions Example

There are three problems in the code



# alarm and pause Functions Example

There are three problems in the code

1. if the caller already has an alarm set, that alarm is erased by the first call to alarm
  - we have to check the return value of alarm, and make sure we wait only until the existing alarm expires



# alarm and pause Functions Example

There are three problems in the code

1. if the caller already has an alarm set, that alarm is erased by the first call to alarm
  - we have to check the return value of alarm, and make sure we wait only until the existing alarm expires
2. disposition for SIGALRM is modified
  - if others are going to use this call, make sure the function is restored after use





# alarm and pause Functions Example

There are three problems in the code

1. if the caller already has an alarm set, that alarm is erased by the first call to alarm
  - we have to check the return value of alarm, and make sure we wait only until the existing alarm expires
2. disposition for SIGALRM is modified
  - if others are going to use this call, make sure the function is restored after use
3. there is race condition between the first call to alarm and the call to pause
  - use setjmp or sigprocmask with sigsuspend



# Example cont'd codes/sleep-pause2.c

to avoid race condition SVR2 used setjmp and longjmp

```
1  #include <setjmp.h>
2  #include <signal.h>
3  #include <unistd.h>
4
5  static jmp_buf env_alrm;
6
7  static void
8  sig_alrm(int signo)
9  {
10     longjmp(env_alrm, 1);
11 }
12
13 unsigned int
14 sleep2(unsigned int seconds)
15 {
16     if (signal(SIGALRM, sig_alrm) == SIG_ERR)
17         return(seconds);
18     if (setjmp(env_alrm) == 0) {
19         alarm(seconds); /* start the timer */
20         pause(); /* next caught signal wakes us up */
21     }
22     return(alarm(0)); /* turn off timer, return unslept time */
23 }
```



## Example cont'd codes/tsleep.c |

There is subtle problem with sleep2 function when it interacts with other signals

In this example, we are trying to make it execute longer than 5 seconds (the argument to sleep2())

```
1  #include <setjmp.h>
2  #include <signal.h>
3  #include <unistd.h>
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <errno.h>
7
8  unsigned int sleep2(unsigned int);
9  static void sig_int(int);
10
11 static jmp_buf env_alrm;
12
13 int
14 main(void)
15 {
16     unsigned int unslept;
17
```



## Example cont'd codes/tsleep.c II

```
18     if (signal(SIGINT, sig_int) == SIG_ERR){
19         fprintf(stderr, "signal(SIGINT) error");
20         exit(1);
21     }
22     unslept = sleep2(5);
23     printf("sleep2 returned: %u\n", unslept);
24     exit(0);
25 }
26
27 static void
28 sig_int(int signo)
29 {
30     int    i, j;
31     volatile int k = 0;
32
33     /*
34      * Tune these loops to run for more than 5 seconds
35      * on whatever system this test program is run.
36      */
37     printf("\nsig_int starting\n");
38     for (i = 0; i < 900000; i++)
39         for (j = 0; j < 10000; j++)
40             k += i * j;
41     printf("sig_int finished\n");
42 }
43
```



## Example cont'd codes/tsleep.c III

```
44 static void
45 sig_alrm(int signo)
46 {
47     longjmp(env_alrm, 1);
48 }
49
50 unsigned int
51 sleep2(unsigned int seconds)
52 {
53     if (signal(SIGALRM, sig_alrm) == SIG_ERR)
54         return(seconds);
55     if (setjmp(env_alrm) == 0) {
56         alarm(seconds); /* start the timer */
57         pause(); /* next caught signal wakes us up */
58     }
59     return(alarm(0)); /* turn off timer, return unslept time */
60 }
```



# Example cont'd codes/tsleep.c

make tsleep

We execute the program by `./tsleep` and interrupt the sleep by typing in the interrupt character

```
James@maker:codes$ time ./tsleep
sleep2 returned: 0
```

```
real 0m5.014s
user 0m0.005s
sys 0m0.007s
```


```
James@maker:codes$ time ./tsleep
^C
sig_int starting
sleep2 returned: 0
```

```
real 0m5.008s
user 0m4.590s
sys 0m0.023s
```

we can see that the `longjmp` from the `sleep2` aborted the other signal handler, `sig_int`, even though it wasn't finished



# LAST WORDS



# Last Words

- Prepare for Exam

